

# How to Create EKS Cluster Using eksctl?

How to Create EKS Cluster Using eksctl? (...)



- You can find the source code for this video in my [GitHub Repo](#).

## Create IAM User for eksctl

First of all, we need to create a user to run eksctl. If you are just getting started, you may create or use an existing user with admin privileges. A better approach would be to create a user with only the permissions needed to perform its functions. For example, eksctl user should only have permissions to create and manage EKS clusters. Let me show you how to create one.

- Go to AWS console and select `IAM` service.

In this example, we're going to create a user with minimum IAM policies needed to run the main use cases of eksctl. These are the ones used to run the integration tests. We will use both managed IAM policies, and we will have to create our own policies as well.

- Let's create IAM policies first. Give it a name `EksAllAccess` and don't forget to replace `account_id`.

**EksAllAccess.json**

```

1  {
2      "Version": "2012-10-17",
3      "Statement": [
4          {
5              "Effect": "Allow",
6              "Action": "eks:*",
7              "Resource": "*"
8          },
9          {
10             "Action": [
11                 "ssm:GetParameter",
12                 "ssm:GetParameters"
13             ],
14             "Resource": [
15                 "arn:aws:ssm:*:<account_id>:parameter/aws/*",
16                 "arn:aws:ssm*:parameter/aws/*"
17             ],
18             "Effect": "Allow"
19         },
20         {
21             "Action": [
22                 "kms>CreateGrant",
23                 "kms>DescribeKey"
24             ],
25             "Resource": "*",
26             "Effect": "Allow"
27         },
28         {
29             "Action": [
30                 "logs>PutRetentionPolicy"
31             ],
32             "Resource": "*",
33             "Effect": "Allow"
34         }
35     ]
36 }

```

- The second one is `IamLimitedAccess` policy.

**IamLimitedAccess.json**

```

1  {
2      "Version": "2012-10-17",
3      "Statement": [
4          {
5              "Effect": "Allow",
6              "Action": [

```

```
7     "iam:CreateInstanceProfile",
8     "iam>DeleteInstanceProfile",
9     "iam:GetInstanceProfile",
10    "iam:RemoveRoleFromInstanceProfile",
11    "iam:GetRole",
12    "iam>CreateRole",
13    "iam>DeleteRole",
14    "iam:AttachRolePolicy",
15    "iam:PutRolePolicy",
16    "iam>ListInstanceProfiles",
17    "iam>AddRoleToInstanceProfile",
18    "iam>ListInstanceProfilesForRole",
19    "iam:PassRole",
20    "iam:DetachRolePolicy",
21    "iam>DeleteRolePolicy",
22    "iam:GetRolePolicy",
23    "iam:GetOpenIDConnectProvider",
24    "iam>CreateOpenIDConnectProvider",
25    "iam>DeleteOpenIDConnectProvider",
26    "iam:TagOpenIDConnectProvider",
27    "iam>ListAttachedRolePolicies",
28    "iam:TagRole"
29 ],
30 "Resource": [
31     "arn:aws:iam::<account_id>:instance-profile/eksctl-*",
32     "arn:aws:iam::<account_id>:role/eksctl-*",
33     "arn:aws:iam::<account_id>:oidc-provider/*",
34     "arn:aws:iam::<account_id>:role/aws-service-role/eks-
35 nodegroup.amazonaws.com/AWSServiceRoleForAmazonEKSNodegroup",
36     "arn:aws:iam::<account_id>:role/eksctl-managed-*"
37 ],
38 },
39 {
40     "Effect": "Allow",
41     "Action": [
42         "iam:GetRole"
43     ],
44     "Resource": [
45         "arn:aws:iam::<account_id>:role/*"
46     ]
47 },
48 {
49     "Effect": "Allow",
50     "Action": [
51         "iam>CreateServiceLinkedRole"
52     ],
53     "Resource": "*",
54     "Condition": {
55         "StringEquals": {
56             "iam:AWSPropertyName": [
57                 "eks.amazonaws.com",
58                 "eks-nodegroup.amazonaws.com",
59                 "eks-fargate.amazonaws.com"
60             ]
61         }
62     }
63 }
```

```

60      ]
61      }
62    }
63  ]
64 }

```

We can attach IAM policies directly to the IAM user or follow the best practices and create an IAM group first.

- Create `EKS` IAM group and attach the following policies:
  - `EksAllAccess` (Customer Managed Policy)
  - `IamLimitedAccess` (Customer Managed Policy)
  - `AmazonEC2FullAccess` (AWS Managed Policy)
  - `AWSCloudFormationFullAccess` (AWS Managed Policy)
- After that, we can create an IAM user. Give it a name `eksctl` and place it in the `EKS` group. Don't forget to download credentials; we will use them to create an AWS profile.

Before proceeding to the next step, make sure that you have AWS CLI installed on your working machine. Depending on the operating system, [you can follow one of those instructions](#) to install the tool.

To run `eksctl` locally, you can export environment variables with your access key and secret or create a profile which is more convenient, in my opinion. Example with environment variables:

```

export AWS_ACCESS_KEY_ID=<access-key>
export AWS_SECRET_ACCESS_KEY=<secret-key>

```

If you just run `aws configure`, it will create a default profile. It's fine, but I don't really want to use that `eksctl` user for anything else rather than creating and managing EKS clusters. You can add a `--profile` option; this will create a named profile. Also, select the region that you want to use, in my case `us-east-1`.

```
aws configure --profile eksctl
```

To verify that everything is configured correctly, run the following command:

```
aws sts get-caller-identity --profile eksctl
```

You should get a json object with your user (example).

```
1  {
2    "UserId": "AIDAWFUREJA5405C55JKB",
3    "Account": "<account_id>",
4    "Arn": "arn:aws:iam:<account_id>:user/eksctl"
5 }
```

Every time you want to execute any `aws` or `eksctl` command, you need to add `--profile eksctl`.

## Create Simple Public EKS Cluster Using eksctl

In this section, we will create a simple public EKS cluster using `eksctl`. If you don't have `eksctl` installed, follow one of those [instructions](#). By public, I mean a Kubernetes cluster with nodes that have **public** IP addresses. You don't need a NAT gateway for that setup.

The simplest way is to run `create` command. It will create VPC as well as an EKS cluster for you. It may be the easiest way but rarely what you want.

```
eksctl create cluster --profile eksctl
```

`eksctl` will create CloudFormation stacks to create EKS cluster and instance groups. In case of an error, you can always inspect the stack itself.

You can customize your cluster by providing the flags to the `eksctl` tool. However, the best approach would be to create a config file. In that case, it's much easier to reproduce the same infrastructure and track the changes in the git.

When the creation is completed, `eksctl` will automatically configure the `kubectl` context. You can immediately access the Kubernetes cluster. Run `kubectl get svc` to get Kubernetes api service from default namespace. You can also run `kubectl get nodes` to get all available Kubernetes workers.

To delete the cluster, run the `delete` command and provide the name of the cluster.

```
eksctl delete cluster --name <name> --profile eksctl
```

## Create Private EKS Cluster Using eksctl

For this example, we will create an EKS cluster with only private nodes. You still will be able to expose services to the internet using a load balancer and ingress. It's just Kubernetes nodes will

not be exposed to the internet, which is rarely allowed by the companies.

- Let's create an `eksctl` config to define the cluster properties. You can find all the possible options [here](#).

### 1-example.yaml

```

1  ---
2  apiVersion: eksctl.io/v1alpha5
3  kind: ClusterConfig
4  metadata:
5    name: private-eks
6    region: us-east-1
7    version: "1.21"
8  availabilityZones:
9    - us-east-1a
10   - us-east-1b
11  managedNodeGroups:
12    - name: general
13      privateNetworking: true
14      instanceType: t3.small
15      desiredCapacity: 1
16    - name: spot
17      privateNetworking: true
18      instanceType: t3.small
19      spot: true
20      desiredCapacity: 1
21      labels:
22        role: spot
23      taints:
24        - key: spot
25          value: "true"
26          effect: NoSchedule

```

- To create this cluster, provide config to the `eksctl`.

```
eksctl create cluster -f 1-example.yaml --profile eksctl
```

- Similar command to delete the cluster.

```
eksctl delete cluster -f 1-example.yaml --profile eksctl
```

## Create EKS Cluster in Existing VPC Using eksctl

Often you already have a VPC, and you want to create an EKS cluster in the same network. You can do it with `eksctl`; you just need to provide a few additional variables.

- Create `main` VPC with the following IPv4 CIDR block `10.0.0.0/16`
- Create `igw` Internet Gateway and attach it to the `main` VPC.
- Create 4 subnets in 2 different availability zones.
  - `private-us-east-1a`, CIDR: `10.0.0.0/18`
  - `private-us-east-1b`, CIDR: `10.0.64.0/18`
  - `public-us-east-1a`, CIDR: `10.0.128.0/18`
  - `public-us-east-1b`, CIDR: `10.0.192.0/18`
- Allocate public IP address for NAT Gateway. Give it a name `nat`.
- Create NAT gateway with `nat` name and place it in one of the public subnets.
- Create 2 routing tables. One for private subnets with a default route (`0.0.0.0/0`) to NAT Gateway, and one for public subnets with a default route(`0.0.0.0/0`) to Internet Gateway.
- Associate subnets with proper routing tables.
- Create a new eksctl config with existing VPC and subnets.

## 2-example.yaml

```

1  ---
2  apiVersion: eksctl.io/v1alpha5
3  kind: ClusterConfig
4  metadata:
5    name: existing-vpc
6    region: us-east-1
7    version: "1.21"
8  vpc:
9    id: vpc-0e88c3c17fb0c7cca
10   subnets:
11     private:
12       us-east-1a:
13         id: subnet-0699d559ea175cac2
14       us-east-1b:
15         id: subnet-049a9d986dc0d9aa5
16     public:
17       us-east-1a:
18         id: subnet-0e4f8a8dda322fac4
19       us-east-1b:
20         id: subnet-0068c08ea23be99dc
21   managedNodeGroups:
22   - name: general
23     privateNetworking: true
24     instanceType: t3.small
25     desiredCapacity: 1

```

- Create a new EKS cluster with a existing VPC.

```
eksctl create cluster -f 2-example.yaml --profile eksctl
```

## IAM Roles for Service Accounts

Amazon EKS supports IAM Roles for Service Accounts (IRSA) that allows cluster operators to map AWS IAM Roles to Kubernetes Service Accounts.

This provides fine-grained permission management for apps that run on EKS and use other AWS services. These could be apps that use S3, any other data services (RDS, MQ, STS, DynamoDB), or Kubernetes components like AWS Load Balancer controller or ExternalDNS.

- Let's create an `AllowListAllMyBuckets` IAM policy to allow list all the buckets in this AWS account.

### AllowListAllMyBuckets.json

```

1  {
2      "Version": "2012-10-17",
3      "Statement": [
4          {
5              "Effect": "Allow",
6              "Action": "s3>ListAllMyBuckets",
7              "Resource": "*"
8          }
9      ]
10 }
```

- Next, create an `eksctl` config with open ID connect enabled.

### 3-example.yaml

```

1  ---
2  apiVersion: eksctl.io/v1alpha5
3  kind: ClusterConfig
4  metadata:
5      name: cluster-irsa
6      region: us-east-1
7  availabilityZones:
8      - us-east-1a
9      - us-east-1b
10  iam:
11      withOIDC: true
12      serviceAccounts:
```

```

13  - metadata:
14    name: foo
15    namespace: staging
16    attachPolicyARNs:
17      - arn:aws:iam:<account_id>:policy/AllowListAllMyBuckets
18    # role has to start with eksctl-
19    roleName: eksctl-list-s3-buckets
20    roleOnly: true
21  - metadata:
22    name: cluster-autoscaler
23    namespace: kube-system
24    wellKnownPolicies:
25      autoScaler: true
26    roleName: eksctl-cluster-autoscaler
27    roleOnly: true
28  managedNodeGroups:
29    - name: general
30    tags:
31      # EC2 tags required for cluster-autoscaler auto-discovery
32      k8s.io/cluster-autoscaler/enabled: "true"
33      k8s.io/cluster-autoscaler/<cluster-name>: "owned"
34    desiredCapacity: 1
35    minSize: 1
36    maxSize: 10

```

- Create a new EKS cluster with a single autoscaling node group.

```
eksctl create cluster -f 3-example.yaml --profile eksctl
```

When the cluster is created, you can go to the AWS IAM service and find additional IAM roles that were created by eksctl.

- First, let's test if we can list S3 buckets from the pod. Create and apply deployment with aws sdk.

### s3.yaml

```

1  ---
2  apiVersion: v1
3  kind: Namespace
4  metadata:
5    name: staging
6  ---
7  apiVersion: v1
8  kind: ServiceAccount
9  metadata:
10   name: foo
11   namespace: staging
12   annotations:

```

```

13     eks.amazonaws.com/role-arn: arn:aws:iam::<account_id>:role/eksctl-list-
14     s3-buckets
15     ---
16   apiVersion: v1
17   kind: Pod
18   metadata:
19     name: aws-cli
20     namespace: staging
21   spec:
22     serviceAccountName: foo
23     containers:
24       - name: aws-cli
25         image: amazon/aws-cli
26         command: [ "/bin/bash", "-c", "--" ]
27         args: [ "while true; do sleep 30; done;" ]
28     tolerations:
29       - operator: Exists
          effect: NoSchedule

```

- Create a Kubernetes service account and a pod.

```
kubectl apply -f k8s/s3.yaml
```

- Then you need to ssh to the pod by using `exec` command.

```
kubectl exec -it aws-cli -n staging -- bash
```

- You can use aws cli to list S3 buckets in your account.

```
aws s3api list-buckets --query "Buckets[ ].Name"
```

## Deploy Cluster Autoscaler to EKS

Create autoscaler deployment and apply it to the cluster.

**cluster-autoscaler.yaml**

```

1   ---
2   apiVersion: v1
3   kind: ServiceAccount
4   metadata:
5     name: cluster-autoscaler
6     namespace: kube-system
7     annotations:
8       eks.amazonaws.com/role-arn: arn:aws:iam::<account_id>:role/eksctl-
9       cluster-autoscaler

```

```

10  ---
11  apiVersion: rbac.authorization.k8s.io/v1
12  kind: ClusterRole
13  metadata:
14    name: cluster-autoscaler
15  rules:
16    - apiGroups: [""]
17      resources: ["events", "endpoints"]
18      verbs: ["create", "patch"]
19    - apiGroups: [""]
20      resources: ["pods/eviction"]
21      verbs: ["create"]
22    - apiGroups: [""]
23      resources: ["pods/status"]
24      verbs: ["update"]
25    - apiGroups: [""]
26      resources: ["endpoints"]
27      resourceNames: ["cluster-autoscaler"]
28      verbs: ["get", "update"]
29    - apiGroups: [""]
30      resources: ["nodes"]
31      verbs: ["watch", "list", "get", "update"]
32    - apiGroups: [""]
33      resources: ["namespaces", "pods", "services", "replicationcontrollers",
34      "persistentvolumeclaims", "persistentvolumes"]
35      verbs: ["watch", "list", "get"]
36    - apiGroups: ["extensions"]
37      resources: ["replicasets", "daemonsets"]
38      verbs: ["watch", "list", "get"]
39    - apiGroups: ["policy"]
40      resources: ["poddisruptionbudgets"]
41      verbs: ["watch", "list"]
42    - apiGroups: ["apps"]
43      resources: ["statefulsets", "replicasets", "daemonsets"]
44      verbs: ["watch", "list", "get"]
45    - apiGroups: ["storage.k8s.io"]
46      resources: ["storageclasses", "csinodes", "csidrivers",
47      "csistoragecapacities"]
48      verbs: ["watch", "list", "get"]
49    - apiGroups: ["batch", "extensions"]
50      resources: ["jobs"]
51      verbs: ["get", "list", "watch", "patch"]
52    - apiGroups: ["coordination.k8s.io"]
53      resources: ["leases"]
54      verbs: ["create"]
55    - apiGroups: ["coordination.k8s.io"]
56      resourceNames: ["cluster-autoscaler"]
57      resources: ["leases"]
58      verbs: ["get", "update"]
59  ---
60  apiVersion: rbac.authorization.k8s.io/v1
61  kind: Role
62  metadata:

```

```
63      name: cluster-autoscaler
64      namespace: kube-system
65  rules:
66    - apiGroups: [""]
67      resources: ["configmaps"]
68      verbs: ["create", "list", "watch"]
69    - apiGroups: [""]
70      resources: ["configmaps"]
71      resourceNames: ["cluster-autoscaler-status", "cluster-autoscaler-
72 priority-expander"]
73      verbs: ["delete", "get", "update", "watch"]
74  ---
75  apiVersion: rbac.authorization.k8s.io/v1
76  kind: ClusterRoleBinding
77  metadata:
78    name: cluster-autoscaler
79  roleRef:
80    apiGroup: rbac.authorization.k8s.io
81    kind: ClusterRole
82    name: cluster-autoscaler
83  subjects:
84    - kind: ServiceAccount
85      name: cluster-autoscaler
86      namespace: kube-system
87  ---
88  apiVersion: rbac.authorization.k8s.io/v1
89  kind: RoleBinding
90  metadata:
91    name: cluster-autoscaler
92    namespace: kube-system
93  roleRef:
94    apiGroup: rbac.authorization.k8s.io
95    kind: Role
96    name: cluster-autoscaler
97  subjects:
98    - kind: ServiceAccount
99      name: cluster-autoscaler
100     namespace: kube-system
101  ---
102  apiVersion: apps/v1
103  kind: Deployment
104  metadata:
105    name: cluster-autoscaler
106    namespace: kube-system
107    labels:
108      app: cluster-autoscaler
109  spec:
110    replicas: 1
111    selector:
112      matchLabels:
113        app: cluster-autoscaler
114    template:
115      metadata:
```

```

116     labels:
117       app: cluster-autoscaler
118   spec:
119     serviceAccountName: cluster-autoscaler
120   containers:
121     - image: k8s.gcr.io/autoscaling/cluster-autoscaler:v1.21.0
122       name: cluster-autoscaler
123       resources:
124         limits:
125           cpu: 100m
126           memory: 600Mi
127         requests:
128           cpu: 100m
129           memory: 600Mi
130       # https://github.com/kubernetes/autoscaler/blob/master/cluster-
131       autoscaler/cloudprovider/aws/README.md
132       command:
133         - ./cluster-autoscaler
134         - --v=4
135         - --stderrthreshold=info
136         - --cloud-provider=aws
137         - --skip-nodes-with-local-storage=false
138         - --expander=least-waste
139         - --node-group-auto-discovery=asg:tag=k8s.io/cluster-
140       autoscaler/enabled,k8s.io/cluster-autoscaler/<cluster-name> # Update cluster
141         - --balance-similar-node-groups
142         - --skip-nodes-with-system-pods=false
143       volumeMounts:
144         - name: ssl-certs
145           mountPath: /etc/ssl/certs/ca-certificates.crt
             readOnly: true
             imagePullPolicy: "Always"
       volumes:
         - name: ssl-certs
           hostPath:
             path: "/etc/ssl/certs/ca-bundle.crt"

```

- After you create all the objects needed to deploy the autoscaler, apply the YAML.

```
kubectl apply -f k8s/cluster-autoscaler.yaml
```

- Check the pod status, and it should be in running state.

```
kubectl get pods -n kube-system
```

- It's a good idea to check logs for any errors.

```
kubectl logs -l app=cluster-autoscaler -n kube-system -f
```

- Let's list all the nodes in the cluster.

```
watch kubectl get nodes
```

- Create a deployment with five replicas to test autoscaling.

#### cluster-autoscaler.yaml

```
1  ---
2  apiVersion: apps/v1
3  kind: Deployment
4  metadata:
5    name: nginx
6  spec:
7    replicas: 5
8    selector:
9      matchLabels:
10        app: nginx
11    template:
12      metadata:
13        labels:
14          app: nginx
15      spec:
16        containers:
17          - name: nginx
18            image: nginx:1.14.2
19            ports:
20              - containerPort: 80
21            resources:
22              requests:
23                memory: 2Gi
24                cpu: 250m
```

```
kubectl apply -f k8s/deployment.yaml
```

#### Clean

- Delete eksctl IAM user
- Delete EKS IAM group
- Delete EksAllAccess IAM policy
- Delete IamLimitedAccess IAM policy
- Delete AllowListAllMyBuckets IAM policy
- Delete EKS eksctl delete cluster -f 3-example.yaml --profile eksctl